

Homework 9: Deterministic Execution

Intro

- ☐ Why would a bug happen only on awsrn but not on your development machine? Why would code that works great for you, occasionally crash for your partner, and always for your TA?

▼ Ans

“——These are a derivative of environmental, time, and randomness sources beyond our control.”

```
int main() {
    int i;
    printf("value of i=%d\n", i);
    printf("address &i=%p\n", &i);
    printf("hash of i=%ld\n", ((uintptr_t)&i) & 127);
    return 0;
}
```

- ☐ Run it a few times using **make undef-compare** and compare the results. Are they the same? Why or why not?

▼ Ans

```
make CFLAGS=-O3 clean undef; ./undef
value of i=32765
address &i=0x7ffd706bb8b4
hash of i=52

make CFLAGS=-O1 clean undef; ./undef
value of i=32766
address &i=0x7ffef79af184
hash of i=4
```

- ☐ Fix the code to define the variable and rerun **make undef-compare** to make sure that it worked.

▼ Ans

```
make CFLAGS=-O3 clean undef; ./undef
value of i=100
address &i=0x7ffd20f06a34
hash of i=52

make CFLAGS=-O1 clean undef; ./undef
value of i=100
address &i=0x7ffc6bd6e4484
hash of i=4
```

- ☐ Run **make undef-noaslr** to run the program a few times without ASLR. Is this program deterministic now?

▼ Ans

```
setarch x86_64 -R ./undef
value of i=100
address &i=0x7fffffff734
hash of i=52

setarch x86_64 -R ./undef
value of i=100
address &i=0x7fffffff734
hash of i=52
```

- ☐ Run **make undef-env** to run the command with differing environments. What is printed now?

▼ Ans

```
setarch x86_64 -R env USER=me ./undef
value of i=100
address &i=0x7fffffff734
hash of i=52

setarch x86_64 -R env USER=professoramarasinghe ./undef
value of i=100
address &i=0x7fffffff724
hash of i=36
```

Checkoff Item 1

▼ Hashtable

slot	ptr	size
0		
1		
2		
3		
4		
5		
6		

- ☐ What do you need to make this program deterministic? Modify the Makefile target for **hashtable1000-good** so that all runs succeed.

```
# XXX modify this target so that 1000 runs all succeed
hashtable1000-good: ./hashtable
    bash -c "for((i=0;i<1000;i++)) ./hashtable; done"
```

▼ Ans

```
bash -c "for((i=0;i<1000;i++)) do setarch x86_64 -R ./hashtable 1 0; done" // always
succeed
```

Checkoff Item 2

- ☐ Modify **hashtable1-bad** target so the program always fails. Show your modified Makefile target. You may find it useful to examine **hashtable.c** to see what system arguments it takes.

```
# XXX modify this target so the program always fails
hashtable1-bad: ./hashtable
    ./hashtable
```

▼ Ans

```
bash -c "for((i=0;i<1;i++)) do setarch x86_64 -R ./hashtable 1 5; done" // always fail
```

Checkoff Item 3

- ☐ What was the bug? What is your fix? Rerun make **hashtable1000** to ensure that your fix always works.

```
void hashtable_insert(void* p, int size) {
    assert(ht.entries < TABLESIZE);
    ht.entries++;

    int s = hash_func(p);
    /* open addressing with linear probing */
    do {
        if (!ht.hashtable[s].ptr) {
            ht.hashtable[s].ptr = p;
            ht.hashtable[s].size = size;
            break;
        }
        /* conflict, look for next item */
        s++;
    } while (1);
}
```

▼ Ans

```
void hashtable_insert(void* p, int size) {
    hashtable_lock();
    assert(ht.entries < TABLESIZE);
    /* ensure atomic, can also use std::atomic<int> */
    ht.entries++;
    int s = hash_func(p);
    /* open addressing with linear probing */
    do {
        if (!ht.hashtable[s].ptr) {
            ht.hashtable[s].ptr = p;
            ht.hashtable[s].size = size;
            break;
        }
        /* conflict, look for next item */
        s++;
    }
```

```

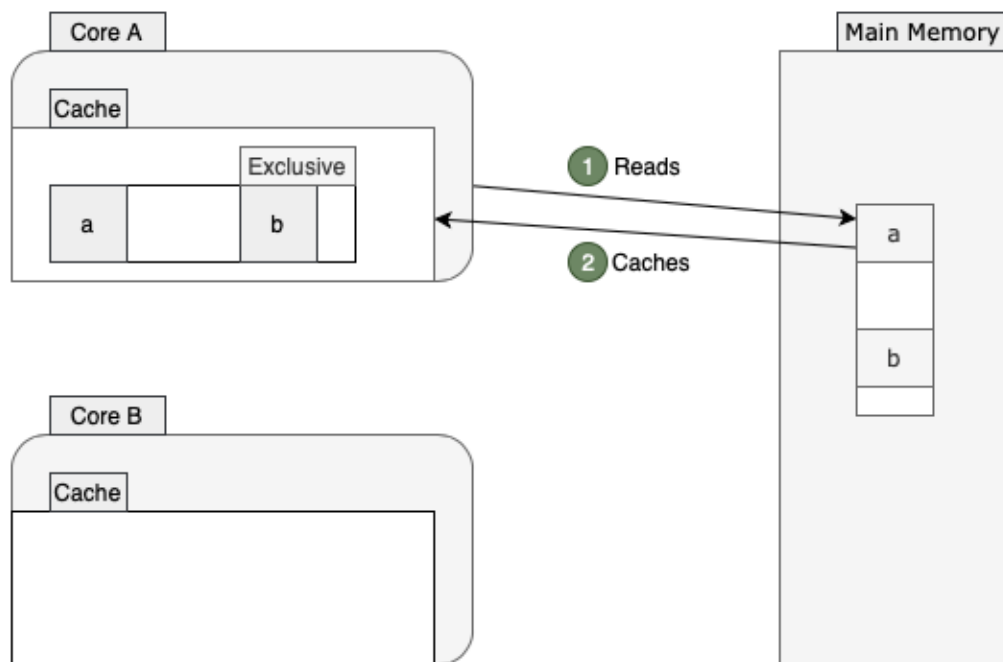
/* ensure the slot can be reset */
s %= TABLESIZE;
} while (1);
hashtable_unlock();
}


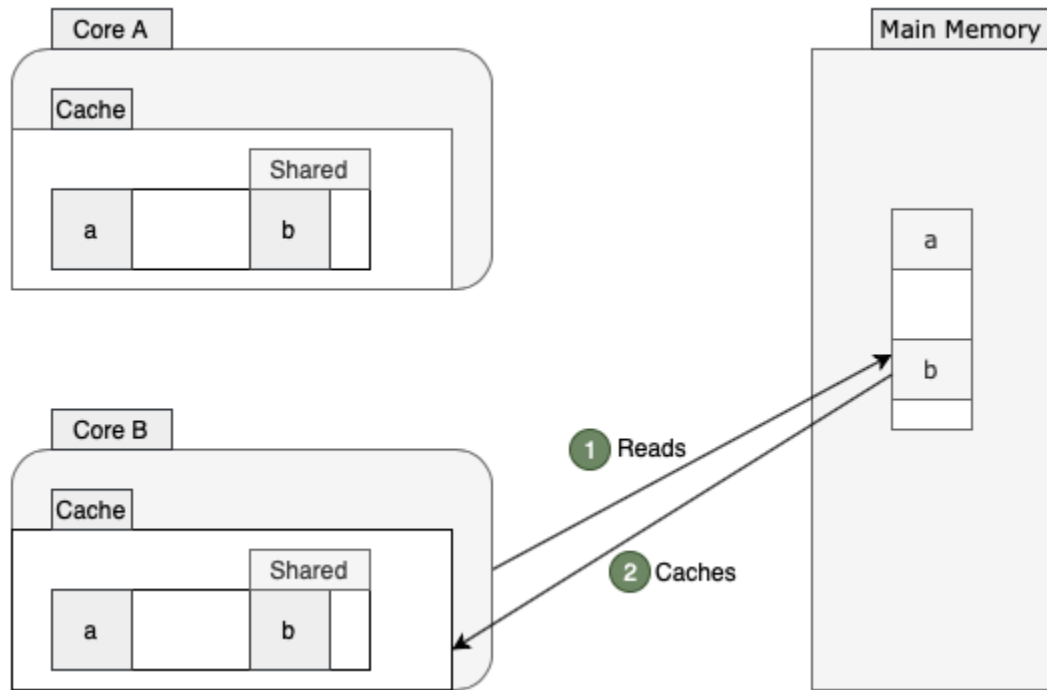
```

Write-up 1

☐ Is the hashlock implementation in **hashlock.c** vulnerable to false sharing? Why or why not?

▼ false sharing





▼ Ans

```
// ori entry
typedef struct entry {
```

```

void* ptr; // 8 bytes
size_t size; // 8 bytes
} entry_t;

// 64 bytes alignas entry
struct alignas(64) entry_t {
    void* ptr;
    size_t size;
};

// ori hashlocks
pthread_mutex_t hashlock[HLOCKS];
#define HASHLOCK_SLOT(l) &hashlock[(l) & (HLOCKS - 1)]

// 64 bytes alignas hashlocks
struct alignas(64) my_lock {
    pthread_mutex_t my_hash_lock;
}
my_lock[HLOCKS];
#define MY_HASHLOCK_SLOT(l) &(my_lock[(l) & (HLOCKS - 1)].my_hash_lock)

// Ans
// In theory yes, but not in practice.

```

Write-up 2

- ☐ What problem is the fairness solution in **hashtable_insert_fair** introducing? Explain or demonstrate the behavior you see.

▼ Ans

hashtable_insert_fair locks and unlocks much more often than **hashtable_insert_locked** when keys of all test cases conflict, which is a huge cost and the actual measured time difference is tens of times larger.

10 threads * 20	all same keys (cost time/lock times/conflict times)	all diff keys (cost time/lock times/conflict times)
hashtable_insert_locked	0.5~3ms / 200 / 19900	0.5~3ms / 200 / 0
hashtable_insert_fair	100ms / 20300 / 19900	0.5~3ms / 400 / 0

Ps: $200 = 10 \text{ threads} * 20$ $20300 = 200 + (1 + 200) * 200 / 2$
 $19900 = (1 + 199) * 199 / 2$

Write-up 3

- ☐ Use `InterlockedCompareExchange64` to implement your changes in **`hashtable_insert_lockless`**. Don't forget to modify **`hashtable_fill`** to use your new code. Run `make hashtable-mt1000` to ensure that it works. What changes are necessary in **`hashtable_lookup`** so you can use just a 8-byte compare-and-swap in **`hashtable_insert_lockless`**?

```
static inline uint64_t InterlockedCompareExchange64(volatile uint64_t* ptr, uint64_t _new,
                                                    uint64_t old) {
    uint64_t prev;
    asm volatile("lock;"
                 "cmpxchgq %1, %2;"
                 : "=a"(prev)
                 : "q"(_new), "m"(*ptr), "a"(old)
                 : "memory");
    return prev;
}
```

▼ CAS(Compare & Swap)

```
uint64_t CAS(uint64_t* ptr, uint64_t _new, uint64_t old) {
    if (*ptr == old) {
        *ptr = _new;
    } else {
        old = *ptr;
    }
    return old;
}
```

▼ Ans

```
void hashtable_insert_lockless(void* p, int size, int tid) {
    hashtable_lock();
    assert(ht.entries < TABLESIZE);
    ht.entries++;
    hashtable_unlock();

    int s = hash_func(p);
    printf("hashtable_insert_lockless: tid:%d slot:%d\n", tid, s);
    /* open addressing with linear probing */
    do {
        if (!InterlockedCompareExchange64((uint64_t*)&ht.hashtable[s].ptr, (uint64_t)p,
0)) {
            ht.hashtable[s].size = size;
            break;
        }
    } while (1);
}
```



```

    }
    s++;
    s %= TABLESIZE;
} while (1);
}

```

Write-up 4

- ☐ (No implementation required.) How would you implement a deterministic hashtable structure, such that the order of insertions does not change the final hashtable state?

▼ Ans

```

#include <iostream>
#include <stdio.h>
#include <typeinfo>
#include <thread>
#include <vector>
#include <mutex>
#include <atomic>
#include <assert.h>

using namespace std;

#define TABLESIZE 256
#define THREADS 10

struct ListNode {
    int ptr;
    int size;
    ListNode* next;
    ListNode() : ptr(0), size(0), next(nullptr) {}
    ListNode(int p, int s) : ptr(p), size(s), next(nullptr) {}
    ListNode(int p, int s, ListNode* n) : ptr(p), size(s), next(n) {}
};

typedef struct hashtable_t {
    ListNode* head_node[TABLESIZE];
    int entries;
    pthread_mutex_t lock;
} hashtable_t;

int hash_func(int ptr) {
    return ptr / 10;
}

void is_greater(int ptr, ListNode* cur_node, bool& res) {
    assert(cur_node);
    res = ptr > cur_node->ptr;
}

```

```

hashtable_t ht = {{}, 0, PTHREAD_MUTEX_INITIALIZER};

void hashtable_lock() {
    pthread_mutex_lock(&ht.lock);
}

void hashtable_unlock() {
    pthread_mutex_unlock(&ht.lock);
}

void hash_insert(int ptr, int size) {
    hashtable_lock();
    assert(ht.entries < TABLESIZE);
    ht.entries++;

    int slot = hash_func(ptr);
    if (ht.head_node[slot] == nullptr) {
        ht.head_node[slot] = new ListNode(ptr, size);
    } else {
        bool res = true;
        ListNode* cur_node = ht.head_node[slot];
        ListNode* pre_node = nullptr;
        while (cur_node) {
            is_greater(ptr, cur_node, res);
            if (!res) break;
            pre_node = cur_node;
            cur_node = cur_node->next;
        }
        ListNode* ptr_node = new ListNode(ptr, size);
        if (res) {
            pre_node->next = ptr_node;
        } else {
            if (pre_node) {
                ptr_node->next = cur_node;
                pre_node->next = ptr_node;
            } else {
                ptr_node->next = cur_node;
                ht.head_node[slot] = ptr_node;
            }
        }
    }
    hashtable_unlock();
}

int main(int argc, char* argv[]) {
    // hash_insert(13, 100);
    // hash_insert(15, 100);
    // hash_insert(11, 100);

    int group = 0;
    if (argc > 1) {
        group = atoi(argv[1]);
    }
}

```

```

vector<shared_ptr<thread>> vec_ptr;
for (int i = 0; i < THREADS; i++) {
    vec_ptr.push_back(make_shared<thread>(hash_insert, group * 10 + i, 100));
}

for (int i = 0; i < THREADS; i++) {
    vec_ptr.at(i)->join();
}

ListNode* cur = ht.head_node[group];
while (cur) {
    cout<<"key:"<<cur->ptr<<" val:"<<cur->size<<endl;
    cur = cur->next;
}

return 0;
}

root@CD-DZ0104843:/home/hanbabang/2_workspace/MIT6_172F18_hw9/MIT6_172F18_hw9# cd "/home/hanbabang/2_workspace/MIT6_172F18_hw9/MIT6_172F18_hw9/" && g++ hashtable_insert_oblivious_.cpp -pthread -o hashtable_insert_oblivious_ && "/home/hanbabang/2_workspace/MIT6_172F18_hw9/MIT6_172F18_hw9/"hashtable_insert_oblivious_ 5
Insert key:50 val:100
Insert key:53 val:100
Insert key:55 val:100
Insert key:52 val:100
Insert key:57 val:100
Insert key:56 val:100
Insert key:54 val:100
Insert key:51 val:100
Insert key:59 val:100
Insert key:58 val:100

Dump    key:50 val:100
Dump    key:51 val:100
Dump    key:52 val:100
Dump    key:53 val:100
Dump    key:54 val:100
Dump    key:55 val:100
Dump    key:56 val:100
Dump    key:57 val:100
Dump    key:58 val:100
Dump    key:59 val:100

```